

1--

```
import pandas as pd
df=pd.read_csv('kc_house_data.csv')
df.head()
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21613 non-null  int64
1   date                 21613 non-null  object
2   price                21613 non-null  float64
3   bedrooms             21613 non-null  int64
4   bathrooms            21613 non-null  float64
5   sqft_living          21613 non-null  int64
6   sqft_lot             21613 non-null  int64
7   floors               21613 non-null  float64
8   waterfront           21613 non-null  int64
9   view                 21613 non-null  int64
10  condition             21613 non-null  int64
11  grade                21613 non-null  int64
12  sqft_above           21613 non-null  int64
13  sqft_basement        21613 non-null  int64
14  yr_built              21613 non-null  int64
15  yr_renovated         21613 non-null  int64
16  zipcode              21613 non-null  int64
17  lat                  21613 non-null  float64
18  long                 21613 non-null  float64
19  sqft_living15        21613 non-null  int64
20  sqft_lot15           21613 non-null  int64
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```

```
df.columns
```

```
Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',
       'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',
       'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',
```

```

        'lat', 'long', 'sqft_living15', 'sqft_lot15'],
        dtype='object')

print(df.isnull().sum())

```

```

id          0
date        0
price       0
bedrooms    0
bathrooms   0
sqft_living 0
sqft_lot    0
floors       0
waterfront  0
view         0
condition    0
grade        0
sqft_above  0
sqft_basement 0
yr_built     0
yr_renovated 0
zipcode      0
lat          0
long         0
sqft_living15 0
sqft_lot15   0
dtype: int64

```

```
print(df.isnull().sum().sum())
```

```
0
```

Le notebook a bien été enregistré.



```

3      9824
4      6882
2      2760
5      1601
6        272
1        199
7         38
8         13
0         13
9          6
10         3
11         1
33         1
Name: bedrooms, dtype: int64

```

```
df['grade'].value_counts()
```

```

7      8981
8      6068
9      2615
6      2038
10     1134
11      399
5       242

```

```

12      90
4       29
13      13
3       3
1       1
Name: grade, dtype: int64


```

```

import matplotlib.pyplot as plt
plt.plot(df['price'])

```


```
[<matplotlib.lines.Line2D at 0x7fc953c53790>]
```




```

import seaborn as sns
sns.distplot(df['price'], bins=5, hist=True, kde=False, color='red')

```

Le notebook a bien été enregistré.  packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is deprecated. Use `displot` instead.

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fc94fc3ae10>
```




```

import matplotlib.pyplot as plt
v=df['grade'].value_counts()
v.plot.bar(rot=45)

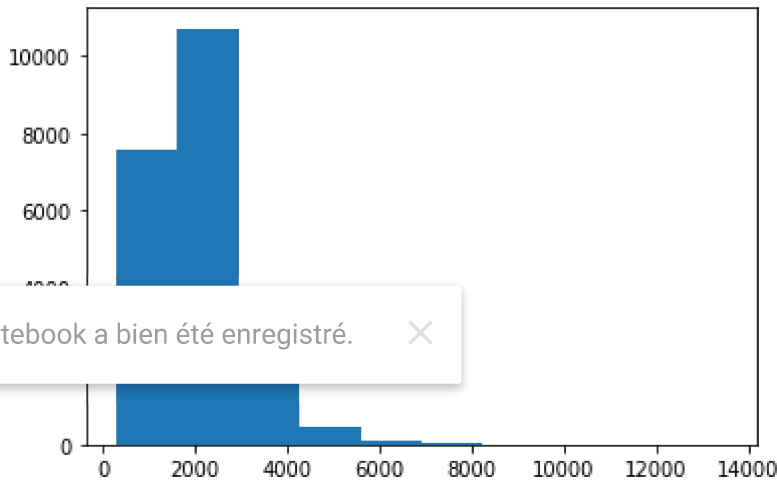
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fc94fa80b10>



```
import matplotlib.pyplot as plt
plt.hist('sqft_living',data=df)
```

```
(array([7.5690e+03, 1.0681e+04, 2.8140e+03, 4.4100e+02, 7.7000e+01,
        2.4000e+01, 2.0000e+00, 3.0000e+00, 1.0000e+00, 1.0000e+00]),
array([ 290., 1615., 2940., 4265., 5590., 6915., 8240., 9565.,
        10890., 12215., 13540.]),
<a list of 10 Patch objects>)
```



Le notebook a bien été enregistré.

3--

```
from sklearn.model_selection import train_test_split
x=df[["id","grade","bathrooms","bedrooms"]]
y=df["price"].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=40)
```

4--

```
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
x=df[["grade"]]
y=df["price"].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=40)
model=LinearRegression()
model.fit(x_train,y_train)
```

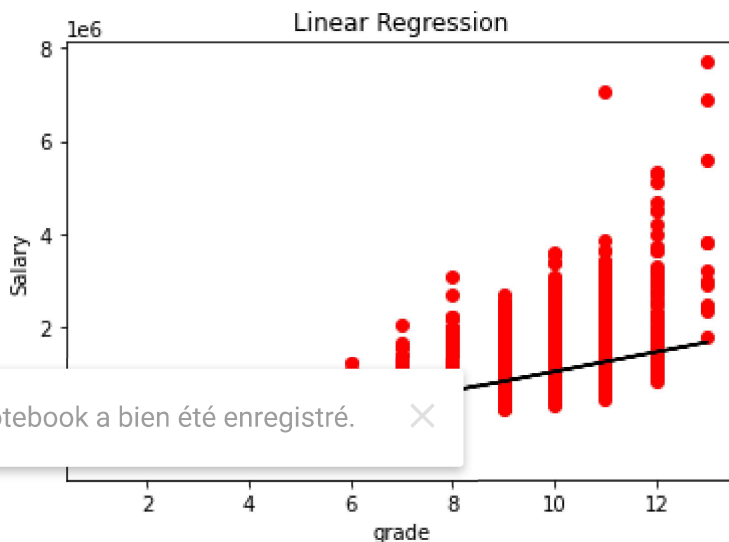
```
predicted=model.predict(x_test)

print("MSE", mean_squared_error(y_test,predicted))
print("R squared", metrics.r2_score(y_test,predicted))

MSE 68533946794.312935
R squared 0.4552042311532961
```

5--

```
plt.scatter(x,y,color="r")
plt.title("Linear Regression")
plt.ylabel("Salary")
plt.xlabel("grade")
plt.plot(x,model.predict(x),color="k")
plt.show()
```



6-- R-squared=0.455<0.5 the half of the output can be explained by the model's inputs donc la correlation est faible

```
x=df[['grade','bedrooms','bathrooms']] #we have more than one input
y=df["price"].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.35,random_state=40) #splitt

model=LinearRegression()
model.fit(x_train,y_train)
predicted=model.predict(x_test)

print("MSE", mean_squared_error(y_test,predicted))
print("R squared", metrics.r2_score(y_test,predicted))

MSE 66370581890.284134
R squared 0.4720658164992009
```

7--

```
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
x =df[['sqft_living','sqft_lot']]
y = df['price'].values
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.35, random_state=40)
lg=LinearRegression()
poly=PolynomialFeatures(degree=2)

x_train_fit = poly.fit_transform(x_train) #transforming our input data
lg.fit(x_train_fit, y_train)
x_test_ = poly.fit_transform(x_test)
predicted = lg.predict(x_test_)

print("MSE: ", metrics.mean_squared_error(y_test, predicted))
print("R squared: ", metrics.r2_score(y_test,predicted))

MSE:  56768005841.851654
R squared:  0.5484479725877804
```

Le notebook a bien été enregistré.



✓ 0 s terminée à 10:42

